

CW1 Part B

The Ultimate Plan

Identifying Bias in Voice-to-Text AI

ECS7025P: Ethics, Regulation and Law in Advanced Digital Information Processing

Finn McKie — Queen Mary University of London

Submission Deadline: 20 March 2026

Narrative Arc: Voice-to-text tools are not simply accent-biased — they are diversity-biased. By comparing real human speakers against AI-generated voices with identical accents, this experiment reveals that the tools are optimised for "clean," standardised speech, systematically disadvantaging the natural variation inherent in real human communication.

How to use this document:

This is your single reference for writing Part B. It contains:

- The exact assignment requirements and marking criteria (with what "Exceptional" looks like)
- Your complete data summary and key findings
- Paragraph-by-paragraph writing plan with guidance on what to include
- Citations with links for further reading
- Lecture references that connect directly to your analysis
- Checklist to ensure you hit every marking criterion

1. ASSIGNMENT REQUIREMENTS

Task: Perform an experiment where 5+ English speakers with different accents use a voice-to-text tool. Compare accuracy, reveal bias, gain insight into training data.

Format:

- 5 to 9 pages (confirmed extended limit)
- 12-point Times or Times New Roman
- Double-spaced
- Saved as PDF
- Name, assignment title, and date in upper left corner (single-spaced)

Required contents:

1. Introduction and identification of 5 speakers
2. Identification of the voice-to-text tool(s) you chose
3. Screenshots of text output for each speaker, captioned with gender and native language
4. Chart of results with clear explanation and interpretation
5. AI tool disclosure statement (or statement that none were used)

■■ DO NOT correct any voice-to-text output. DO NOT accept device corrections. Screenshots must show raw, unedited output. The marking criteria specifically checks for this.

AI Usage Policy (Part B):

You MAY use AI tools to help format and present your data (e.g. chart creation). You MUST disclose all AI tools used and their purpose. The voice-to-text tool itself is acknowledged as AI. Keep in mind: "you must also engage your human brain to ensure the data are represented accurately." The marking criteria rewards human voice and penalises generic chatbot output.

2. MARKING CRITERIA — TARGETING EXCEPTIONAL (70%+)

Each criterion below shows what you need to hit the top bracket. Your plan is designed to exceed every one.

Experimental Design & Data Collection (20%)

■ **Exceptional:** "Methodical selection of 5 or more distinct accents. Screenshots are clear, captioned correctly, and adhere strictly to the no correction rule."

Your status: ■ EXCEEDING. You have 10 real speakers across 10 distinct accents, plus 4 AI voices. This is double the requirement. For the write-up, present your primary 5 speakers (Group A) in the main body, then use Group B and Group C as extended analysis. This shows methodical, deliberate experimental design far beyond what's expected.

Data Visualisation & Accuracy (20%)

■ **Exceptional:** "Data is creatively translated into a compelling visual that is clearly labelled and interpreted. The lecturer is inspired to share it with the class."

Your plan: Two charts minimum. (1) A grouped bar chart comparing accuracy across speakers and tools. (2) A comparison of human vs AI voice accuracy — this is the visual that'll make the lecturer want to share it. Consider a heatmap showing which specific WORDS failed across speakers — reveals patterns at a glance.

Critical Analysis of Bias (30%) — THE BIG ONE

■ **Exceptional:** "Deep exploration of WHY specific errors occurred (phonetics, training data gaps). Connects findings to systemic bias. Goes above and beyond basic assignment."

Your plan: This is where you win. Three layers of analysis:

Layer 1 — Phonetic patterns: Identify specific L1 transfer effects (e.g. "th"→"t/d" for Nigerian/Indian speakers, "pl"→"bl/pr" for Arabic/Mandarin speakers, vowel shifts for Scottish). Everyone will do this — it's your baseline.

Layer 2 — Training data hypothesis: Use your two-tool comparison and AI voice data to argue that the models are optimised for standardised broadcast speech. Cite Koenecke et al. (2020).

Layer 3 — Systemic implications: Connect to algorithmic bias (Week 6), abstracted power (Week 4), and Kant's Formula of Humanity. This is the layer that puts you above everyone else.

Structure & Academic Writing (10%)

■ **Exceptional:** "Effectively uses 5-9 pages (extended limit confirmed). Writing is compelling, eloquent, and uses a high-level academic tone. Human voice is apparent, as are a few human errors."

Key: Write it yourself. Don't over-polish. A few natural spelling or grammar imperfections are BETTER than robotic perfection — the criteria explicitly says so. Be eloquent but be YOU.

AI Tool Disclosure & Integrity (10%)

■ **Exceptional:** "Transparent, detailed acknowledgment of AI assistance OR a clear statement explaining that AI was not used."

Your plan: Be fully transparent. Disclose: (1) Apple Dictation and QuillBot as the experimental tools, (2) ElevenLabs for generating AI accent voices (Group C), (3) Any charting tools used. Transparency here is a strength, not a weakness — especially since your ElevenLabs experiment IS the innovative element.

3. YOUR DATA AT A GLANCE

Reference sentence:

"Please call Stella, ask her to bring six spoons of fresh snow peas, five thick slabs of blue cheese, and a small plastic snake for the kids."

Speaker Profiles

Speaker	Accent	Native Language	Gender	Group
Austin	US English	English	M	A
Finn (you)	British English	English	M	A
Hugh	Scottish	English	M	A
Megan	Irish	English	F	A
Yomi	Nigerian English	Nigerian	M	A
Ansh	Indian	Hindi	M	B
Ella	German	German	F	B
Aaron	Chinese	Mandarin	M	B
Mahmoud	Arabic	Arabic	M	B
Alex	French	French	M	B

Accuracy Summary — Apple Dictation

Speaker	Accent	Errors	Accuracy	Key Errors
Austin	US English	1	96.3%	"thick" → "six"
Finn	British	4	85.2%	"snake"→"strength", "kids"→"kit", omissions
Hugh	Scottish	3	88.9%	"ella" inserted, "snow peas"→"Snoopy's"
Megan	Irish	4	85.2%	"Ella" inserted, "snow"→"no", "five thick"→"56"
Yomi	Nigerian	2	92.6%	"snow peas"→"snobby"
Ansh	Indian	2	92.6%	"snow peas"→"nappies"
Ella	German	2	92.6%	"thick"→"tick", "the"→"ze"
Aaron	Chinese	4	85.2%	"bring"→"bling", "fresh"→"flesh", "plastic"→"prastic"
Mahmoud	Arabic	3	88.9%	"snow peas"→"Snoopy's", "thick"→"fix"
Alex	French	0	100%	Perfect transcription

Accuracy Summary — QuillBot

Speaker	Accent	Errors	Accuracy	Key Errors
Austin	US English	0	100%	Perfect transcription
Finn	British	1	96.3%	"small" omitted
Hugh	Scottish	1	96.3%	"snow"→"shnow"
Megan	Irish	3	88.9%	"snow"→"new", "thick"→"fix", "snake"→"plake"
Yomi	Nigerian	3	88.9%	"ask"→"aks", "thick"→"tick", "the"→"de"
Ansh	Indian	3	88.9%	"five"→"fiwe", "thick"→"tick", "the"→"de"
Ella	German	2	92.6%	"thick"→"tick", "the"→"ze"
Aaron	Chinese	5	81.5%	"bring"→"bling", "fresh"→"flesh", "blue"→"boo", etc.
Mahmoud	Arabic	4	85.2%	"Please"→"Blease", "spoons"→"sboons", etc.
Alex	French	3	88.9%	"her"→"er", "thick"→"sick", "the"→"ze"

AI Voices — Apple Dictation (Group C)

Voice	Accent	Errors	Accuracy	Notes
ElevenLabs	British (AI)	0	100%	Perfect transcription
ElevenLabs	Indian (AI)	1	96.3%	"spoons"→"phones"
ElevenLabs	French (AI)	0	100%	Perfect transcription
ElevenLabs	Scottish (AI)	1	96.3%	"spoons"→"bins"

KEY FINDINGS:

1. AI voices averaged 98.1% accuracy vs human speakers ~89-91%. Same accents, dramatically different results. The tools are optimised for "clean" standardised speech, not real human diversity.
2. Apple was better for non-native speakers (91.9%) than native (89.6%). QuillBot was the opposite — native (94.1%) vs non-native (87.4%). The two tools have **inverted bias profiles**, proving bias is tool-specific and training-data-specific, not an inherent limitation of speech recognition.

4. KEY PATTERNS IN YOUR DATA

Pattern 1: The "th" problem

The word "thick" was misrecognised for 7 out of 10 speakers on at least one tool. The dental fricative /θ/ does not exist in most world languages. German, French, Nigerian English, Hindi, Mandarin, and Arabic speakers all produced substitutions: "tick", "fix", "sick", "six". This is the most systematic error pattern in your data — it reveals the tool's expectation of a specific phoneme that most English speakers worldwide do not produce the "standard" way.

Pattern 2: "snow peas" — the compound noun problem

This phrase was mangled in wildly different ways: "Snoopy's" (Scottish, Arabic on Apple), "snobby" (Nigerian), "nappies" (Indian), "shnow" (Scottish on QuillBot). The tools clearly lack context for this less-common compound noun when pronunciation varies even slightly. Interesting contrast: all AI voices got "snow peas" correct — because they pronounced it with textbook clarity.

Pattern 3: Consonant cluster substitution (L1 transfer)

Aaron (Mandarin L1): "bring"→"bling", "fresh"→"flesh", "plastic"→"prastic" — consistent /r/→// substitution. This is a well-documented phonological feature of Mandarin-influenced English where /r/ and // are not contrastive phonemes. The tool faithfully transcribed what it heard, but the PATTERN reveals it cannot contextually correct for known L1 transfer effects.

Mahmoud (Arabic L1): "Please"→"Blease", "spoons"→"sboons", "peas"→"beas", "plastic"→"blastic" — systematic /p/→/b/ substitution. Arabic lacks the voiceless bilabial plosive /p/, so speakers often produce /b/. Again, the tool transcribes literally rather than inferring contextually.

Pattern 4: Tool architecture differences

Apple Dictation produced substitutions, omissions AND insertions (added "ella" for Scottish/Irish speakers). QuillBot produced almost exclusively substitutions with zero insertions. This suggests fundamentally different model approaches: Apple attempts to "fill in" uncertain audio (generative), while QuillBot maps more directly (discriminative). Both fail, but in characteristically different ways — worth a paragraph in your analysis.

Pattern 5: Tool-specific bias inversion

Your group averages reveal something striking: **Apple was better for non-native speakers (Group B: 91.9%) than native speakers (Group A: 89.6%). QuillBot was the opposite — better for native (94.1%) than non-native (87.4%).** The two tools have inverted bias profiles. This means bias isn't universal — it's tool-specific and training-data-specific. This is a powerful argument against anyone who claims "accents are just hard for computers." No — different tools fail differently for different people, which proves it's a design and data problem, not an inherent technical limitation.

Pattern 6: The Finn paradox — your own data as evidence

You (Finn) scored 85.2% on Apple Dictation — one of the worst in Group A, despite being a native UK English speaker. This is counterintuitive: Apple's training data should favour you. This proves that single-attempt accuracy is volatile — background noise, speaking speed, slight mumbling can all affect one reading. But here's the crucial argument: **multiple retries would likely improve everyone's scores, but that's not how people actually use voice-to-text.** Nobody dictates a message three times. This experiment deliberately mirrors real-world conditions. Your low score doesn't undermine the findings — it reinforces them. Even the "favoured" group isn't immune to errors, but the PATTERN across all speakers still clearly shows bias disproportionately affects certain accents.

Pattern 7: The French paradox

Alex (French) scored 100% on Apple Dictation — the only human to achieve perfect transcription. Yet on QuillBot he scored 88.9% (3 errors). Meanwhile Aaron (Chinese) scored poorly on BOTH tools. This inconsistency between tools for the same speaker reinforces that bias is tool-specific and training-data-specific, not inherent to the accent itself.

5. LECTURE CONNECTIONS — STAND OUT FROM THE CLASS

These are direct connections between YOUR data and Dr Peterson's teaching. Weaving these in shows you've been paying attention and thinking critically — lecturers notice this.

Week 6: Algorithmic Bias (4 March 2026)

Dr Peterson's definition: "The influence of human prejudice on an algorithm's training data, design, embedded assumptions, or end-usage that leads to unfair or discriminatory outcomes."

Her voice-to-text example: "My iPhone dictation has never correctly spelled any of my Indian in-laws' or friends' names, even with repeated use."

Your connection: Your data shows this isn't anecdotal — it's systematic. Ansh (Indian accent) had "snow peas" → "nappies" on Apple, and "five" → "fiwe", "the" → "de" on QuillBot. The tool Dr Peterson described failing on Indian names is the SAME tool (Apple Dictation) that struggled with your Indian speaker. Reference her directly.

Week 6: Kantian Ethics — Formula of Humanity

Kant's principle: "You should always treat others as ends-in-themselves and never as mere means to an end."

Dr Peterson's AI update: "If you intend to treat them/their data as mere means to an end, respect their autonomy by giving them a choice to avoid being used in this way."

Your argument: When voice-to-text tools work brilliantly for American/RP English speakers but systematically fail for Nigerian, Indian, Arabic, and Chinese speakers, those users aren't being treated as "ends" — people the tool was designed to serve equally. They're afterthoughts. The tool was optimised for one group and everyone else must adapt. This is using diverse users as "mere means" — their data may have been collected to improve the tool, but the tool wasn't improved to serve them.

Week 4: Abstracted Power (20 Feb 2026)

Dr Peterson's concept: "Abstracted power is a human actor's influence or control over a system, process, or dataset which, as a function of the technology that enables it, obscures or distances the human actor from consequences of that influence or control."

Citation: Peterson, T.L., Ferreira, R. and Vardi, M.Y. (2023) "Abstracted Power and Responsibility in Computer Science Ethics Education," IEEE Transactions on Technology and Society.

Your argument: The engineers who curate speech training data exercise abstracted power. They decide which accents, dialects, and speech patterns are "standard" — invisible decisions that determine who the tool serves well and who it fails. Your AI voice experiment makes this visible: when the same accent is produced "cleanly" by ElevenLabs, the tool works perfectly. The bias isn't in the accent — it's in the gap between "ideal" and "real" speech, a gap that training data curation choices created.

Week 1: Automation Bias

Dr Peterson defined automation bias as "when we over-rely on automated aids and decision support systems, or become complacent in assuming the technology is always correct." A brief mention: if users trust voice-to-text output without checking, accent-based errors could propagate into documents, messages, and records — compounding the harm.

6. CITATIONS & FURTHER READING

Use 4-6 citations maximum. This is a personal experiment, not a literature review. Each citation should serve a specific analytical purpose.

Citation	Why Use It	Link
Koenecke, A. et al. (2020) "Racial disparities in automated speech recognition" <i>Proceedings of the National Academy of Sciences</i> , 117(14), 7684-7689.	THE benchmark paper. Found ASR error rates were nearly double for Black speakers (35%) vs white speakers (19%) across 5 major commercial systems. Directly supports your findings.	https://www.pnas.org/doi/10.1073/pnas.1915768...
Tatman, R. (2017) "Gender and Dialect Bias in YouTube's Automatic Captions" <i>Proceedings of the First ACL Workshop on Ethics in NLP</i> .	Found systematic accuracy differences by dialect and gender in Google's ASR. Scottish English and speakers from New Zealand/South Africa were least accurate. Directly relevant to your Hugh (Scottish) data.	https://aclanthology.org/W17-1606/...
Peterson, T.L., Ferreira, R. and Vardi, M.Y. (2023) "Abstracted Power and Responsibility in Computer Science Ethics Education" <i>IEEE Transactions on Technology and Society</i> .	Your own lecturer's published research on abstracted power. Citing your lecturer's peer-reviewed work in context shows genuine engagement.	https://ieeexplore.ieee.org/document/10125029...
Harwell, D. (2018) "The accent gap: How Amazon's and Google's smart speakers leave certain voices behind" <i>The Washington Post</i> .	Accessible journalism piece documenting the same phenomenon you're studying. Good for framing the real-world impact beyond academic settings.	https://www.washingtonpost.com/technology/201...
Markl, N. (2022) "Language variation and NLP: Bias is a feature, not a bug" <i>Proceedings of the 2nd Workshop on Trustworthy NLP (TrustNLP)</i> .	Argues that ASR bias against non-standard varieties isn't a technical flaw but a design choice. Powerful framing for your systemic analysis.	https://aclanthology.org/2022.trustnlp-1.11/...

Lecture references (not formal citations, but weave in naturally):

- Dr Peterson, Week 6 lecture — algorithmic bias definition, voice-to-text example, Kantian ethics
- Dr Peterson, Week 4 lecture — abstracted power concept (cite her published paper above for formal reference)
- Dr Peterson, Week 1 lecture — automation bias (brief mention in your analysis)

7. PARAGRAPH-BY-PARAGRAPH WRITING PLAN

This is your writing blueprint. Each section tells you exactly what to cover. YOU write it in your own voice — don't let anyone or anything write it for you. The marking criteria rewards imperfect human writing over polished chatbot output.

■■ REMEMBER: "Human voice is apparent, as are a few human errors" = Exceptional (70%+). "Text is generic chatbot-style output, verbose and free of errors" = Satisfactory (50-59%). Write like YOU, not like a machine.

SECTION 1: Introduction (~0.5 page)

Paragraph 1: Context and tool choice

- Open with WHY this experiment matters — voice-to-text is everywhere (phones, accessibility, transcription)
- Name your tools: Apple Dictation (iPhone native) and QuillBot
- Briefly explain why you chose two tools — to compare whether bias is tool-specific or universal
- Mention you also tested AI-generated voices as a control group (tease the innovation)
- State the reference sentence

SECTION 2: Speaker Profiles (~0.5 page)

Paragraph 2: Introduce your speakers

- Present all 10 speakers with accent, gender, native language
- Explain your selection rationale — you deliberately chose accents from different language families
- Group A (5 native/near-native English), Group B (5 non-native English), Group C (4 AI voices)
- A simple table works well here — captioned properly

Paragraph 3: Methodology note

- All speakers read the sentence in person, at natural speed, into the same device
- No corrections accepted, raw screenshots taken
- Explain the ElevenLabs setup briefly: you generated accented AI voices and played them to the tool
- **Methodology limitation (2-3 sentences):** Each speaker read the sentence once with no retries. This mirrors real-world usage but means individual results could be influenced by momentary factors like speaking speed or background noise. Use YOUR OWN result as evidence: you scored 85.2% on Apple despite being a native UK speaker — proof that single attempts are volatile. But crucially, the overall pattern of bias still holds across all speakers, which makes the findings MORE robust, not less.

SECTION 3: Screenshots (~1–1.5 pages)

- Include screenshots for your primary 5 speakers (Group A) — required by assignment
- Caption each: "Speaker: [Name], Gender: [M/F], Native Language: [X], Tool: [Y]"
- You can include selected Group B/C screenshots if space allows, or reference them in appendix

- Make sure screenshots are CLEAR and READABLE — resize if needed

SECTION 4: Data Visualisation & Interpretation (~1 page)

Chart 1: Accuracy comparison across all speakers

- Grouped bar chart: each speaker on x-axis, two bars (Apple vs QuillBot), y-axis = accuracy %
- Colour-code by accent group (native English vs non-native vs AI)
- This immediately shows the pattern: AI voices cluster near 100%, humans spread widely

Chart 2 (optional but recommended): Human vs AI accuracy

- Simple comparison: average human accuracy vs average AI voice accuracy
- This is your "share with the class" visual — the gap is striking

Interpretation paragraph:

- Walk through the charts — which speakers did the tool favour? (Austin/Alex on Apple, Austin on QuillBot)
- Which speakers were most disadvantaged? (Aaron consistently worst, Mahmoud on QuillBot)
- Note the French paradox: Alex = 100% on Apple but 88.9% on QuillBot. Tool-specific bias.

SECTION 5: Critical Analysis of Bias (~1.5–2 pages) — WHERE YOU WIN

Paragraph 6: Phonetic patterns (your baseline)

- "thick" as the most-failed word — explain /θ/ as a rare phoneme globally
- Aaron's /r/→/l/ pattern (Mandarin L1 transfer)
- Mahmoud's /p/→/b/ pattern (Arabic lacks /p/)
- Cite Tatman (2017) on dialect-specific error patterns

Paragraph 7: Training data hypothesis (your differentiator)

- **The Apple argument:** Apple is a US company — founded, headquartered, and predominantly staffed in the US. Their Dictation training data was almost certainly dominated by US English speech samples — collected from US users, tested by US QA teams, optimised for US accent patterns. Austin's top score (96.3% Apple, 100% QuillBot) is not coincidence — the tool works best for the people closest to the people who built it. That's textbook algorithmic bias.
- **The QuillBot contrast:** QuillBot is cloud-based, potentially trained on more diverse web-sourced data, which could explain its different bias profile — better for native English overall but with different failure patterns.
- American English speaker (Austin) performed best on BOTH tools
- AI voices (clear, standardised pronunciation) achieved near-perfect accuracy
- Argue: training data likely dominated by American/RP broadcast-quality speech
- Cite Koenecke et al. (2020): commercial ASR systems showed nearly double error rates for Black speakers
- Your finding is consistent with their conclusion: "these tools have likely been over-trained on data that favours certain demographics"

Paragraph 8: The AI voice revelation (your unique contribution)

- This is what nobody else will have — lead with it
- AI British voice = 100% vs real British speaker (Finn) = 85.2%. Same accent, 15% gap.
- AI Scottish voice = 96.3% vs real Scottish (Hugh) = 88.9%.
- Argument: the bias is not against accents per se, but against the NATURAL VARIATION in real speech — prosody, rhythm, hesitation, coarticulation
- Cite Markl (2022): "Language variation and NLP: Bias is a feature, not a bug"
- This reframes the entire finding: the tools are designed for idealised speech, not real people

Paragraph 8.5: The Finn paradox — single-attempt volatility

- Your own result (85.2% Apple) as a native UK speaker — counterintuitive and powerful
- Acknowledge: multiple retries would improve scores for everyone
- But: real-world usage IS single-attempt. Nobody dictates three times. Your methodology is deliberate.
- Key line: "The fact that even a native speaker in the tool's expected demographic can score poorly underscores that reliability is inconsistent for everyone — but that inconsistency disproportionately affects certain accents."

Paragraph 8.75: Tool-specific bias inversion (bonus insight)

- Apple performed better for non-native speakers; QuillBot performed better for native speakers
- This proves bias is not inherent to accents — it's a product of each tool's training data
- Brief but powerful point: if bias were simply "accents are hard," both tools would fail the same way

Paragraph 9: Systemic implications — connect to lectures (your killer section)

- Reference Dr Peterson's Week 6 definition of algorithmic bias and her voice-to-text example
- Connect to abstracted power (Week 4, Peterson et al. 2023): training data curation as invisible power
- Apply Kant's Formula of Humanity: diverse speakers treated as afterthoughts, not ends-in-themselves
- Brief mention of automation bias (Week 1): if users trust the output, errors propagate
- End with: who bears the cost? Those already marginalised by language barriers now face an additional digital barrier

SECTION 6: AI Tool Disclosure (~0.25 page)

- State clearly: Apple Dictation and QuillBot used as experimental voice-to-text tools
- ElevenLabs used to generate AI accent voices for the control group experiment
- State any charting/visualisation tools used (e.g. Excel, Python matplotlib)
- If you used AI for anything else (e.g. Atlas helping plan structure), disclose it
- Be transparent and specific — this is worth 10% and "transparent, detailed acknowledgment" = Exceptional

8. PRE-SUBMISSION CHECKLIST

FORMAT

- 5-9 pages (double-spaced, 12pt Times/Times New Roman) — extended limit confirmed
- Name, assignment title, date in upper left (single-spaced)
- Saved as PDF
- Titled per QM Plus submission directions

EXPERIMENTAL DESIGN (20%)

- 5+ distinct accents identified with clear descriptions
- All screenshots present, captioned with gender + native language
- NO corrections made to any output (the "no correction" rule)
- Tool(s) clearly identified

DATA VISUALISATION (20%)

- At least one chart present
- Chart is clearly labelled with axes, legend, title
- Data interpretation paragraph accompanies the chart
- Visual is creative enough to "inspire the lecturer to share it"

CRITICAL ANALYSIS (30%)

- Specific phonetic error patterns identified and explained
- Training data hypothesis supported with evidence
- Connection to systemic/societal bias (not just technical errors)
- References to at least 2 academic sources
- Goes "above and beyond" the basic assignment

ACADEMIC WRITING (10%)

- Human voice apparent throughout
- Academic tone maintained
- Writing flows logically, not just listing results
- A few natural imperfections (NOT zero errors)

AI DISCLOSURE (10%)

- All AI tools identified with specific purposes
- Transparent and detailed

■ Includes voice-to-text tools + any additional tools

Now go write it. Your data is excellent, your angle is unique, and your lecture connections will put you above the rest. Write in YOUR voice — imperfect, thoughtful, and genuine. Come back when you've drafted each section and we'll review together. Let's get that A.